

# Linear Regression and Gradient Descent

**Prof. Mingkui Tan**

SCUT Machine Intelligence Laboratory (SMIL)



# Contents

- 1 Introduction to Machine Learning
- 2 Linear Regression
- 3 Closed-form Solution
- 4 Gradient Descent

# Contents

**1** Introduction to Machine Learning

**2** Linear Regression

**3** Closed-form Solution

**4** Gradient Descent

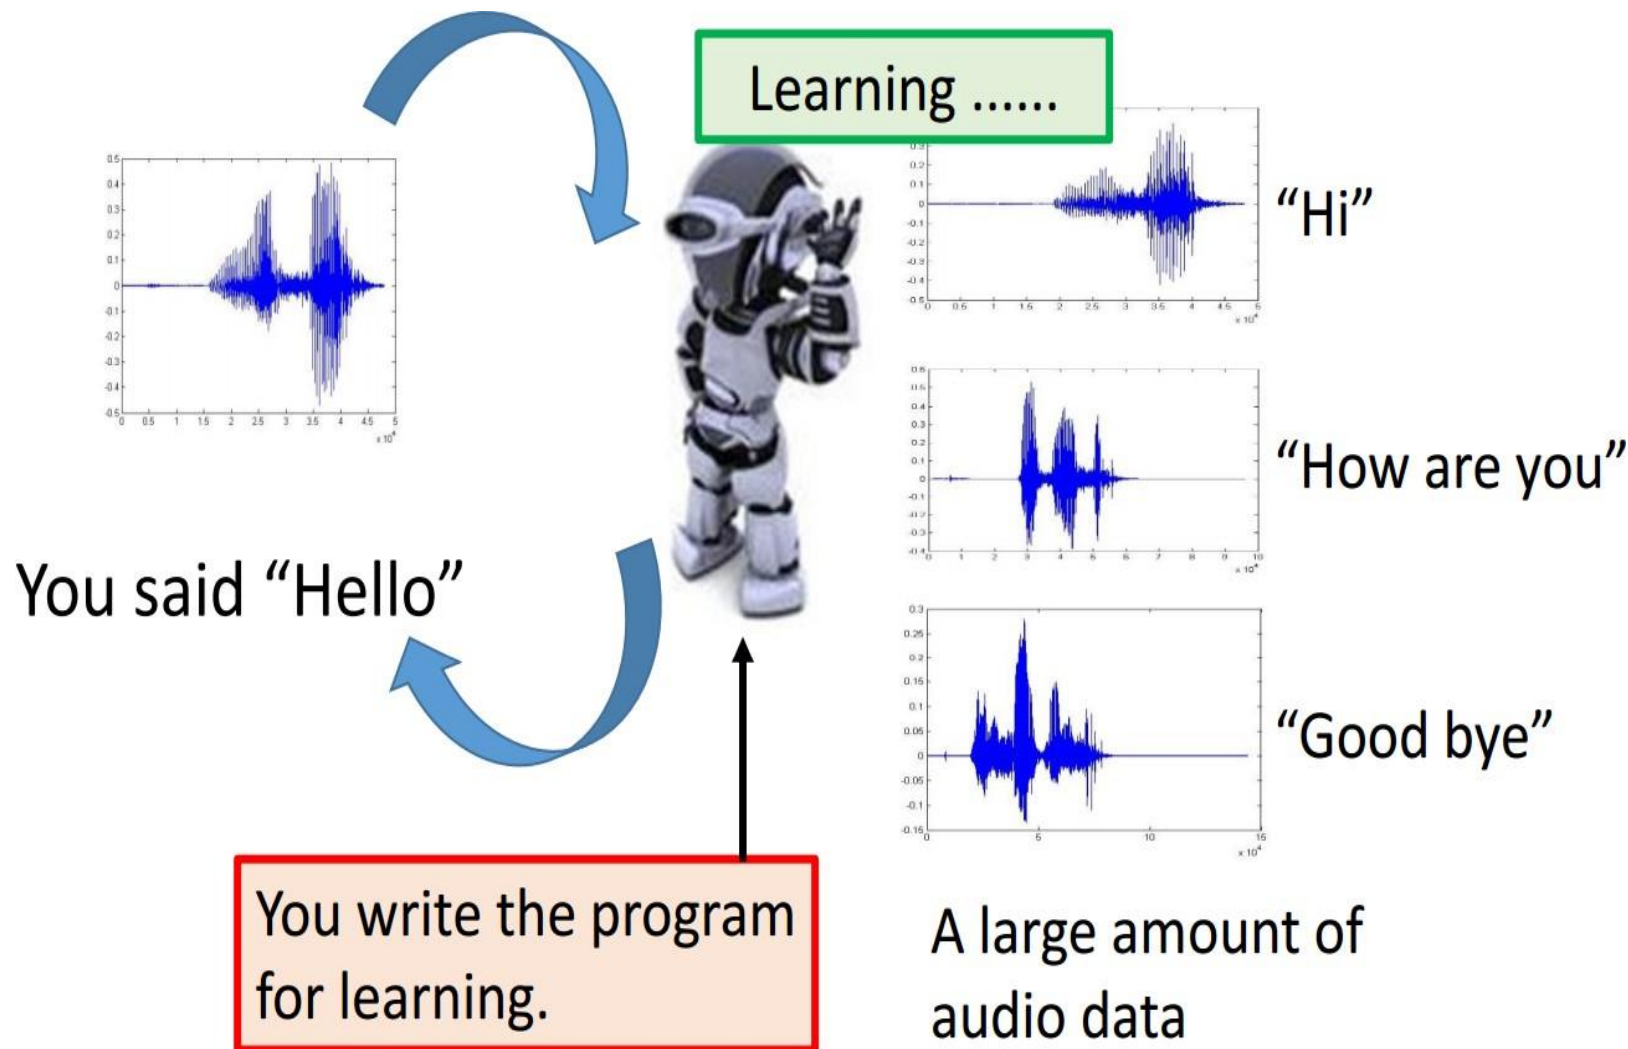
# Introduction to Machine Learning

## What is Machine Learning?

Machine Learning composes of three parts:

- Data
- Model
- Loss Function

# Introduction to Machine Learning



Speech Recognition

# Introduction to Machine Learning

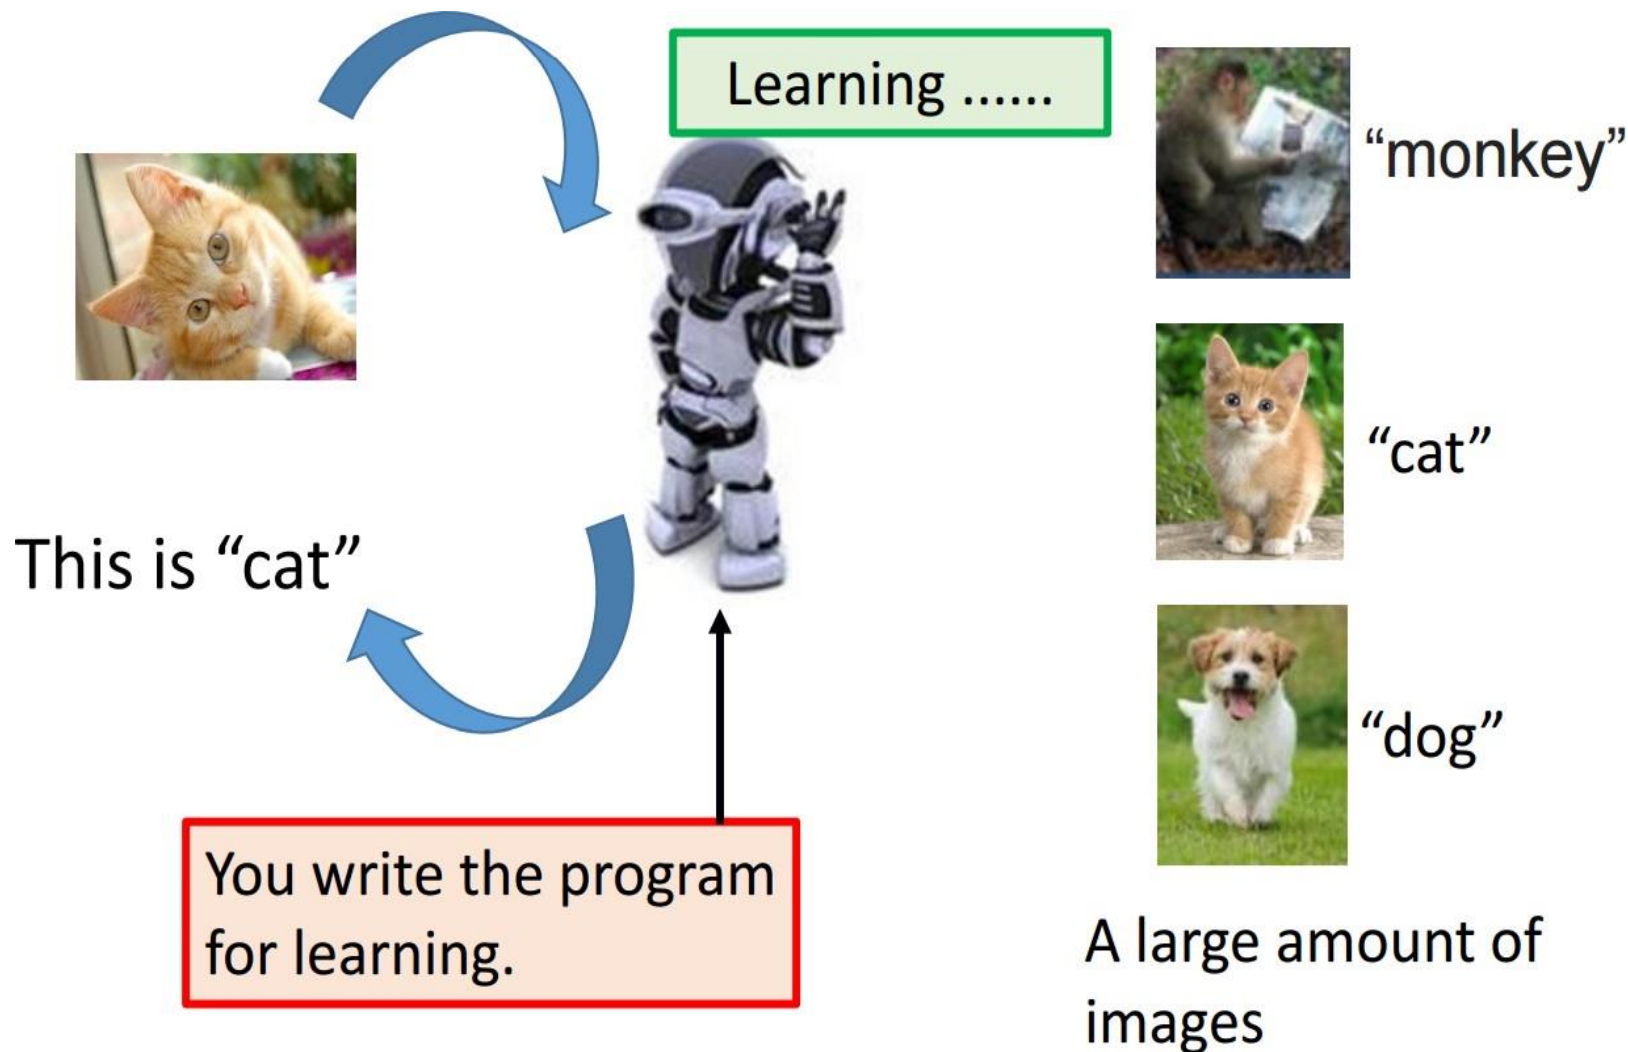
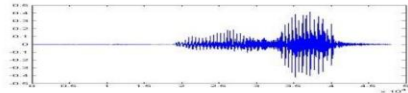


Image Recognition

# Introduction to Machine Learning

## Machine Learning $\approx$ Looking for a Function


### ■ Speech Recognition

$$f(\text{  }) = \text{"How are you"}$$

### ■ Image Recognition

$$f(\text{  }) = \text{"Cat"}$$

### ■ Playing Go

$$f(\text{  }) = \text{"5-5"}_{\text{(next move)}}$$

### ■ Dialogue System

$$f(\text{ "Hi" }) = \text{ "Hello" }$$

(what the user said)      (system response)

# Introduction to Machine Learning

A set of  
function

Model

$f_1, f_2 \dots$

$f_1$



) = "cat"

$f_2$



) = "money"

$f_1$



) = "dog"

$f_2$

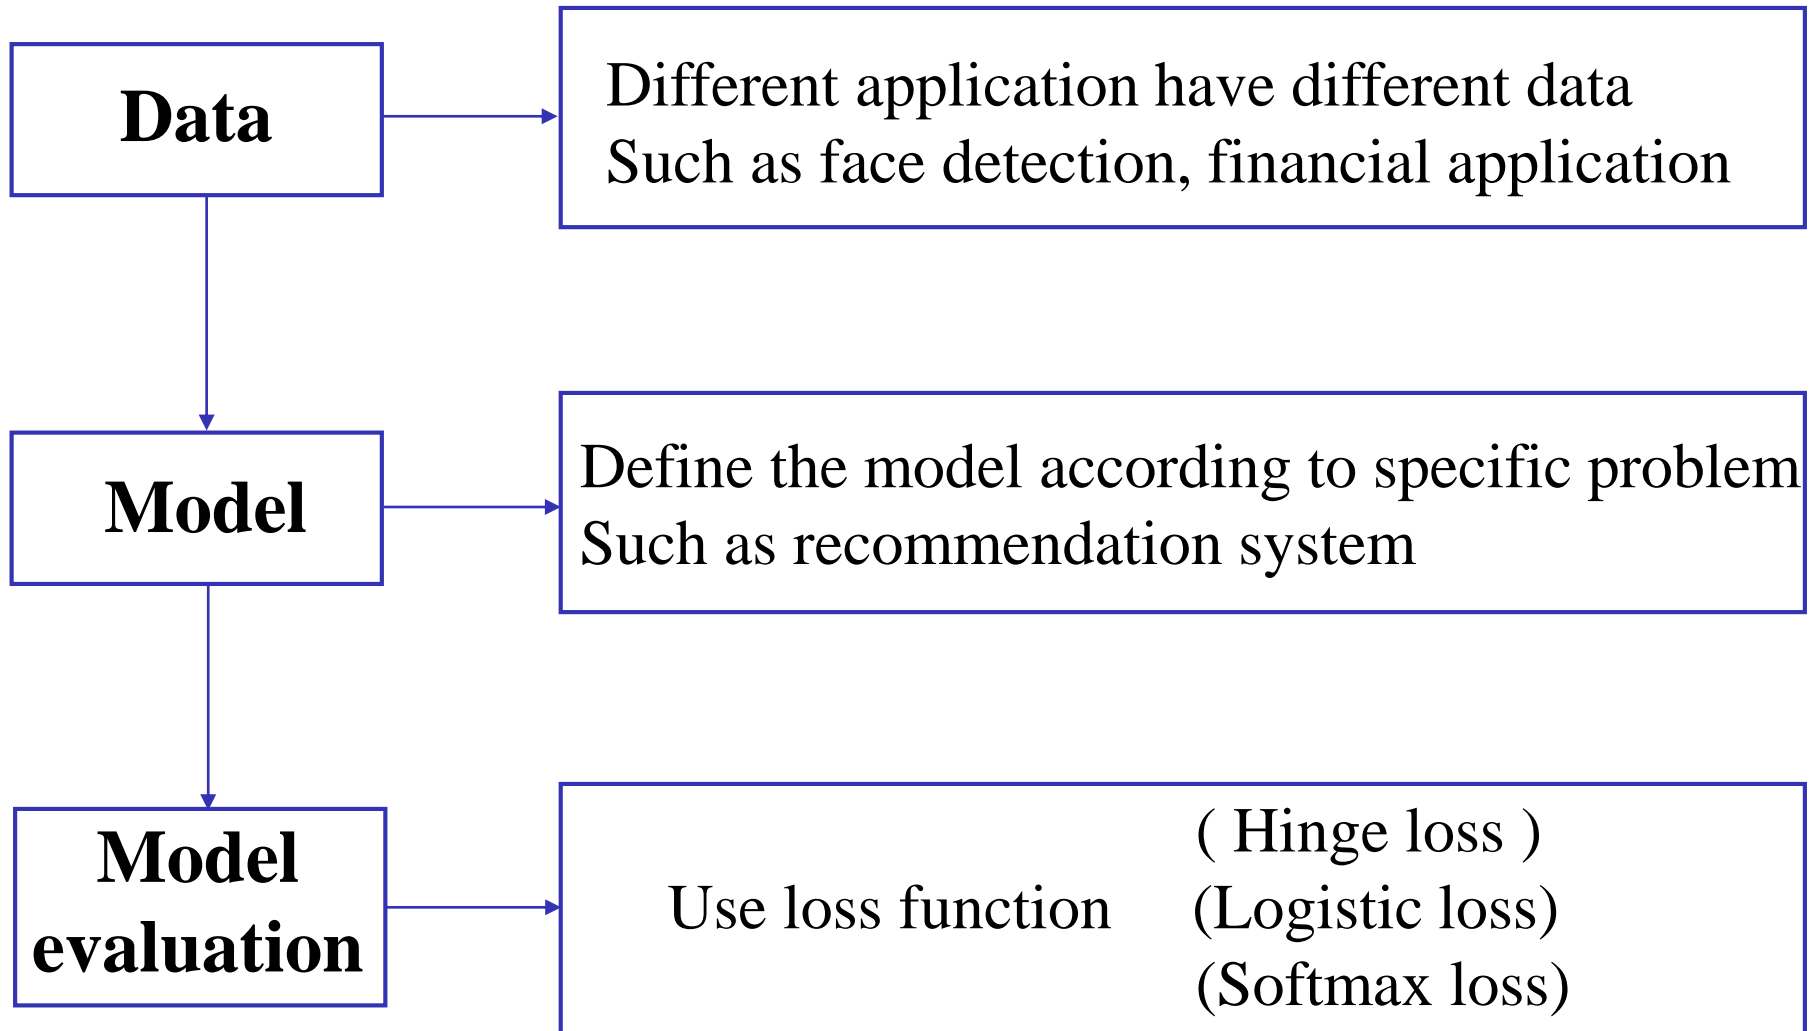


) = "snake"

Image Recognition



# Three Main Elements of Machine Learning




# Introduction to Machine Learning

**Machine Learning is so simple...**



Just like putting an elephant into the fridge...



# Column Vector

Data:

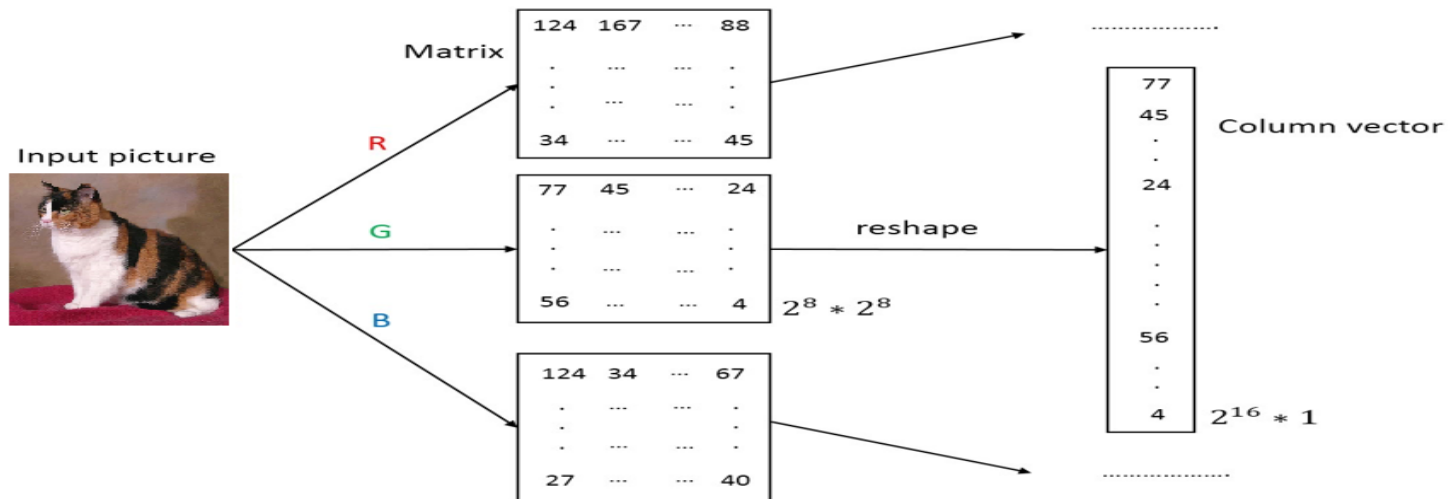
$$\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$$

$\mathbf{x}$  is the input, which is usually presented as a **column vector**

$y$  is the output, for example, a person's name

$n$  is the number of samples

For example,  $\mathbf{x}$  can be a picture stored as a matrix:



# Introduction to Machine Learning

- Use a function to predict  $y$ :

$$\hat{y} = f(x)$$

- However, the prediction may be inconsistent with the ground-truth
- Calculate the difference by loss function:

$$\mathcal{L}_{\mathcal{D}}(\mathbf{W}) = \sum_{i=1}^n l(\hat{y}_i, y_i)$$

where  $\mathcal{D}$  refers to data and  $\mathbf{W}$  refers to parameter

# Regression

Loss:

■ Absolute value loss:

$$l(\hat{y}_i, y_i) = |\hat{y}_i - y_i|$$

■ Least squares loss:

$$l(\hat{y}_i, y_i) = \frac{1}{2} (\hat{y}_i - y_i)^2$$

Total loss function:

$$\mathcal{L}_{\mathcal{D}}(\mathbf{W}) = \sum_{i=1}^n l(\hat{y}_i, y_i)$$

# Regression

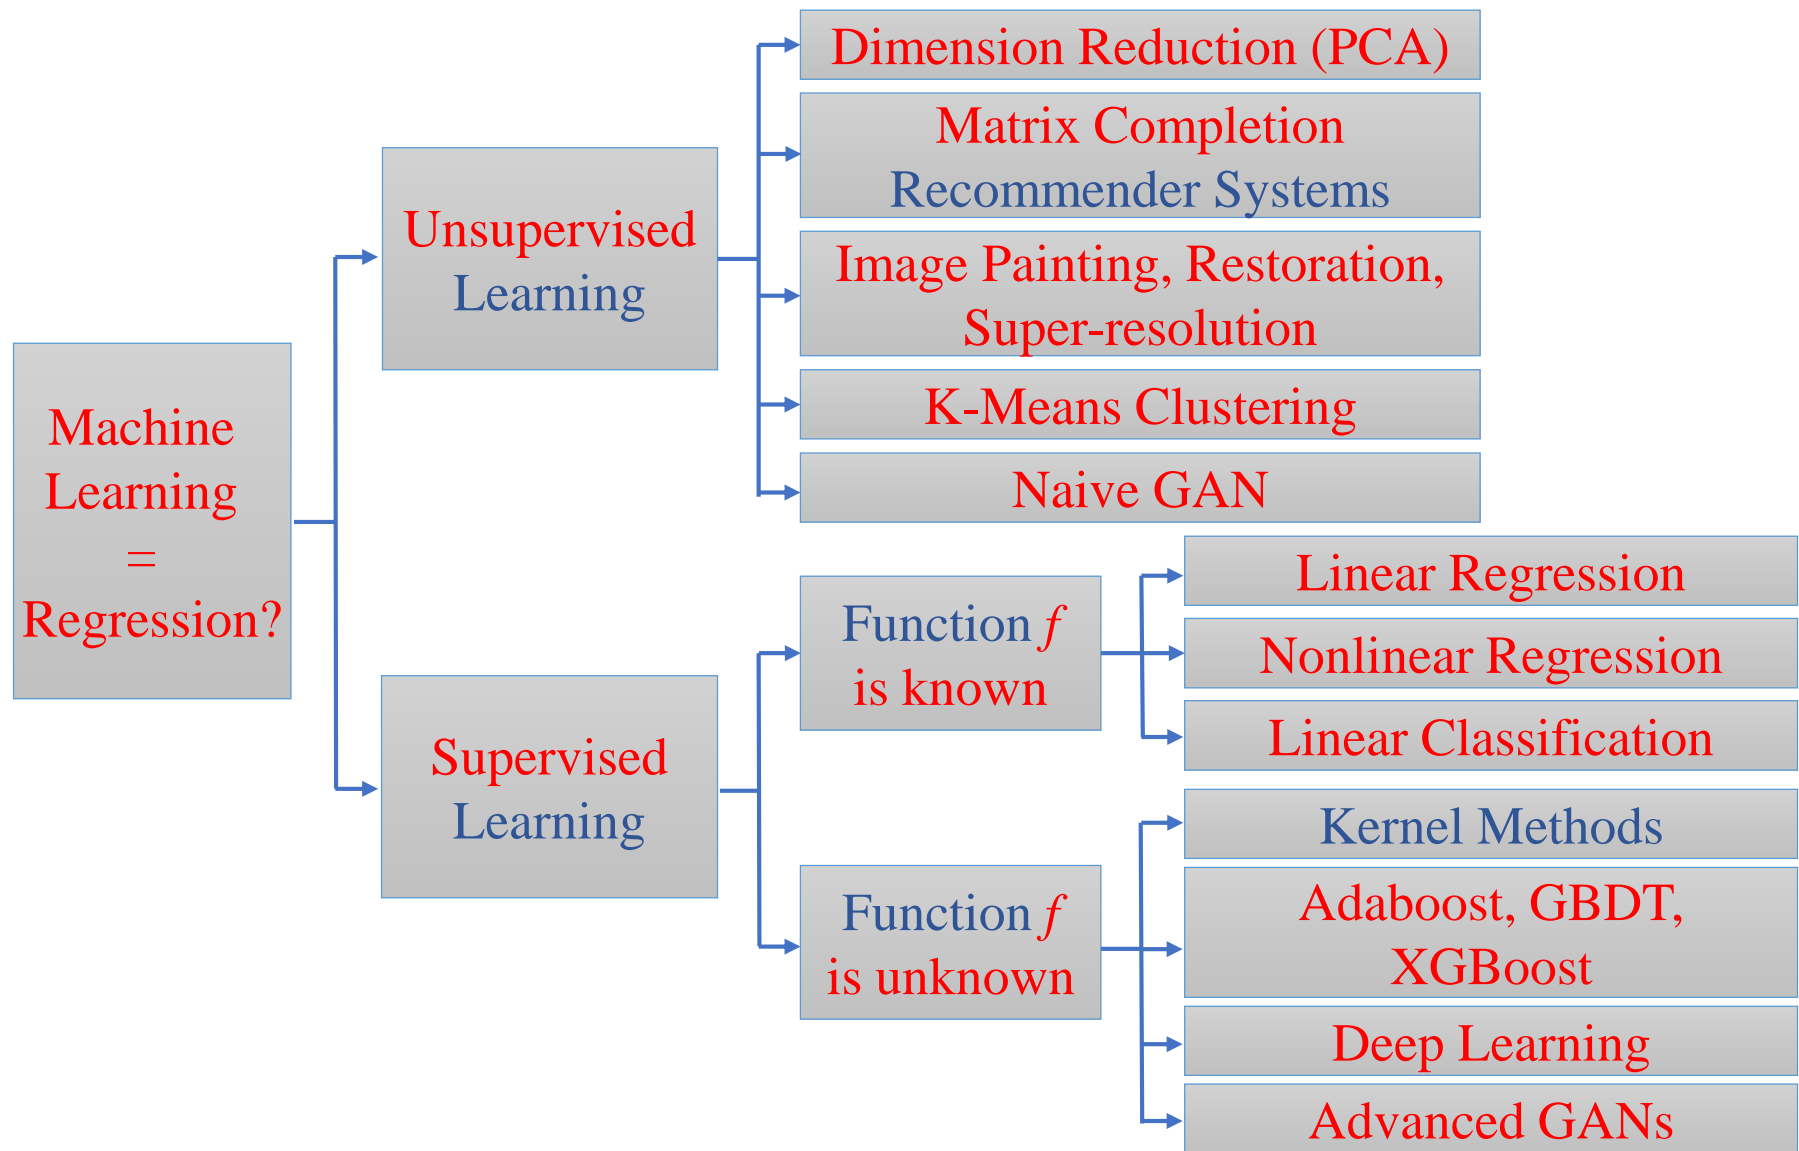
- The smaller value of  $\mathcal{L}_{\mathcal{D}}$  is better, and loss function  $\mathcal{L}_{\mathcal{D}}$  plays a major role in machine learning

## Target:

- Find the best  $f$  by solving the following optimization problem:

$$f^* = \operatorname{argmin}_f \sum_{i=1}^n l(f(x), y_i)$$

# Machine Learning



# Supervised Machine Learning

■ Supervised learning is the machine learning task of inferring a function from **labeled training data**

Labeled data





cat



dog

Unlabeled data





# Dataset for Supervised Learning

## Libsvm dataset

- It contains many classification, regression, multi-label and string data sets

<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/>

- You can use LIBSVM, a package, with these sets

<http://www.csie.ntu.edu.tw/~cjlin/libsvm>

- You can also use LIBLINEAR, a linear classifier, with the sets

<https://www.csie.ntu.edu.tw/~cjlin/liblinear/#document>

- Other tutorials you can read are as follows:

Tools: <https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/>

Guide: <https://www.csie.ntu.edu.tw/~cjlin/papers/guide/guide.pdf>

# Introduction to the Format of LIBSVM

## Two properties of data:

- The number of features is large
- Each instance is sparse for most feature values are zero

## Sparse format:

<label1> <index1>:<value1> <index2>:<value2> ...

<label2> <index1>:<value1> <index2>:<value2> ...

- An example for classification:

+1 1:2 4:5 \n -1 2:4 \n

translate to: The points (2,0,0,5) and (0,4,0,0) are assigned to class +1 and class -1 respectively

# Contents

1 Introduction to Machine Learning


2 Linear Regression

3 Closed-form Solution

4 Gradient Descent

# Linear Regression

**Simple linear regression** describes the linear relationship between a variable  $x$  and a response variable  $y$



Which  
one is  
better?

Simple linear 1-D regression

# Problem Setup for Regression

## ■ Inputs

Input space:  $\mathcal{X} = \{\mathbf{x}_i\}_{i=1}^N, \mathbf{x}_i \in \mathbb{R}^m$

$N$  is the number of data samples

$\mathbf{x}_i$  includes  $m$  features

## ■ Outputs

Output space:  $\mathcal{Y} = \{y_i\}_{i=1}^N, y_i \in \mathbb{R}$

## ■ Goal

Learn a hypothesis / model  $f: \mathcal{X} \rightarrow \mathcal{Y}$

# Linear Regression


Learn  $f(\mathbf{x}; \mathbf{w}, b)$  with

- Parameters:  $\mathbf{w} \in \mathbb{R}^m, b \in \mathbb{R}$
- Input:  $\mathbf{x}$  where  $x_i \in \mathbb{R}$ , features for  $i \in \{1, \dots, m\}$
- Model Function:

$$\begin{aligned} f(\mathbf{x}; \mathbf{w}, b) &= w_1 x_1 + \dots + w_m x_m + b \\ &= \sum_{i=1}^m w_i x_i + b \\ &= \mathbf{w}^T \mathbf{x} + b \end{aligned}$$

# Linear Regression

## ■ What makes a good model?



# Performance Measure for Regression

## ■ Least squared loss

$$\begin{aligned}\mathcal{L}_{\mathcal{D}}(\mathbf{w}, b) &= \frac{1}{2} \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \mathbf{w}, b))^2 \\ &= \frac{1}{2} \sum_{i=1}^n (y_i - \hat{y}_i)^2\end{aligned}$$

Training: find minimizer of least squared loss

$$\mathbf{w}^*, b^* = \operatorname{argmin}_{\mathbf{w}, b} \mathcal{L}_{\mathcal{D}}(\mathbf{w}, b)$$



# Contents

1 Introduction to Machine Learning

2 Linear Regression

3 Closed-form Solution

4 Gradient Descent

# Matrix Presentation for Loss Function

In order to simplify our proof, we introduce augmented matrix and augmented vector and still represent them by  $\mathbf{w}$  and  $\mathbf{X}$ .

$$\mathbf{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_i, \dots, \mathbf{x}_n)^T$$

i.e.

$$\mathbf{x}_i = (1, x_{i1}, x_{i2}, \dots, x_{in})$$

$$\mathbf{w} = (b, w_1, w_2, \dots, w_n)^T$$

Loss function:

$$\mathcal{L}_{\mathcal{D}}(\mathbf{w}) = \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$

$$\text{where } \mathbf{X} = \begin{pmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_n \end{pmatrix} = \begin{pmatrix} 1 & x_{11} & \cdots & x_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{nn} \end{pmatrix}, \mathbf{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}$$

# Matrix Presentation for Loss Function

■ Proof:

$$\begin{aligned}\mathcal{L}_D(\mathbf{w}) &= \frac{1}{2} \sum_{i=1}^n (y_i - \mathbf{x}_i^T \mathbf{w})^2 \\&= \frac{1}{2} \begin{bmatrix} y_1 - \mathbf{x}_1^T \mathbf{w} \\ \vdots \\ y_n - \mathbf{x}_n^T \mathbf{w} \end{bmatrix}^T \begin{bmatrix} y_1 - \mathbf{x}_1^T \mathbf{w} \\ \vdots \\ y_n - \mathbf{x}_n^T \mathbf{w} \end{bmatrix} \\&= \frac{1}{2} \left( \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix} \mathbf{w} \right)^T \left( \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix} \mathbf{w} \right) \\&= \frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \\&= \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2\end{aligned}$$

# Analytical Solution

How to address the linear regression question?

■ Closed-form solution to linear regression:

$$\mathcal{L}_D(\mathbf{w}) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}), \text{ Let } \mathbf{a} = \mathbf{y} - \mathbf{X}\mathbf{w},$$

$$\begin{aligned} \frac{\partial \mathcal{L}_D(\mathbf{w})}{\partial \mathbf{w}} &= \frac{\partial \mathbf{a}}{\partial \mathbf{w}} \frac{\partial (\frac{1}{2} \mathbf{a}^T \mathbf{a})}{\partial \mathbf{a}} \\ &= \frac{1}{2} \frac{\partial \mathbf{a}}{\partial \mathbf{w}} (2\mathbf{a}) \\ &= \frac{\partial (\mathbf{y} - \mathbf{X}\mathbf{w})}{\partial \mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w}) \\ &= -\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \end{aligned}$$

Since  $\mathcal{L}_D(\mathbf{w})$  is a convex function,  $\frac{\partial \mathcal{L}_D(\mathbf{w})}{\partial \mathbf{w}} = 0$  derive  $\mathbf{w}^*$

# Analytical Solution

- Assuming  $|\mathbf{X}^T \mathbf{X}| \neq 0$

- Let 
$$\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} = -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{w} = 0$$

$$\Rightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\Rightarrow \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- Solve the optimal parameter  $\mathbf{w}^*$

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}_{\mathcal{D}}(\mathbf{w}) = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

# Challenges about Analytical Solution

There are two challenges left to address about the analytical solution  $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$  :

- Many matrices are not invertible

Necessary and Sufficient Condition:

If  $\mathbf{X}$  is a matrix of  $m$  rows and  $n$  columns ( $n \leq m$ ),

$$|\mathbf{X}^T \mathbf{X}| \neq 0 \iff \text{rank}(\mathbf{X}) = n$$

- The inverse of a large matrix needs huge memory, which takes  $O(m^3)$  to compute.

# Contents

1 Introduction to Machine Learning

2 Linear Regression

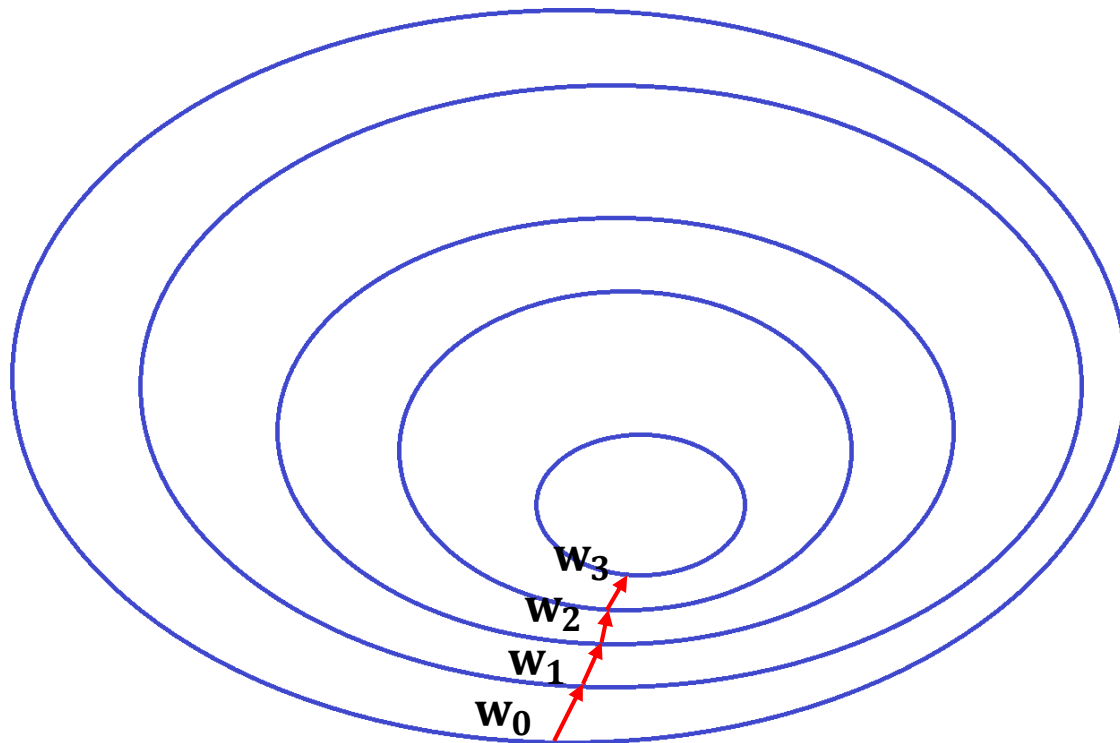
3 Closed-form Solution

4 Gradient Descent

# Gradient Descent

- Get the best  $\mathbf{w}$  by minimizing a loss function  $\mathcal{L}_{\mathcal{D}}(\mathbf{w})$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}_{\mathcal{D}}(\mathbf{w})$$





# Descent Direction

- We use  $\mathbf{d} = -\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$  as the direction of optimization
- Gradient (vector of partial derivatives)

$$\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial w_1} \\ \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial w_m} \end{bmatrix}$$

(We always write a vector into column form)

- Why  $\mathcal{L}_{\mathcal{D}}(\mathbf{w}') = \mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) \leq \mathcal{L}_{\mathcal{D}}(\mathbf{w}), \eta \rightarrow 0^+ ?$

# Descent Direction

Proof:

By Taylor expansion, when  $\eta \rightarrow 0^+$ :

$$\begin{aligned}\mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) &= \mathcal{L}_{\mathcal{D}}(\mathbf{w}) + \left( \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \eta \mathbf{d} + o(\eta \mathbf{d}) \\ &= \mathcal{L}_{\mathcal{D}}(\mathbf{w}) + \eta' \left( \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \mathbf{d}\end{aligned}$$

Note that  $\eta' > 0$  and

$$\eta' \left( \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \mathbf{d} = -\eta' \mathbf{d}^T \mathbf{d} \leq 0$$

We have:


$$\mathcal{L}_{\mathcal{D}}(\mathbf{w}') = \mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) \leq \mathcal{L}_{\mathcal{D}}(\mathbf{w})$$

# Gradient Descent: Update Parameters

Minimize loss by repeated gradient steps (when no closed form):

- Compute gradient of loss with respect to parameters  $\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$
- Update parameters with learning rate  $\eta$

$$\mathbf{w}' = \mathbf{w} - \eta \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$$



# Appropriate Value of Learning Rate

**Learning rate  $\eta$**  has a large impact on convergence

- Too large  $\eta \Rightarrow$  oscillate and may even diverge
- Too small  $\eta \Rightarrow$  too slow to converge

Adaptive learning rate (For example) :

- Set larger learning rate at the beginning
- Use relatively smaller learning rate in the later epochs
- Decrease the learning rate:

$$\eta^{t+1} = \frac{\eta^t}{t + 1}$$

Thank You