



Big Data and IBM Solution

Zeming Zhao

zhaozm@cn.ibm.com

November 1, 2014

© 2014 IBM Corporation

Agenda

- Big Data
- Hadoop Brief Introduction
- HDFS Structure and Characteristic
- MapReduce Model
- Q&A

What is Big Data



- Big data is being generated by everything **around us** at all times. Every digital process and social media exchange produces it. Systems, sensors and mobile devices transmit it.
- Big data is an all-encompassing term for any collection of data sets so **large** and **complex** that it becomes difficult to process using traditional data processing applications.

3Vs Characteristics of Big Data

- Volume: How much data
- Velocity: How fast data is processed
- Variety: The various types of data

Changes made by Big Data

- **Competitive advantage**

Data is emerging as the world's newest resource for competitive advantage.



- **Decision making**

Decision making is moving from the elite few to the empowered many.



- **Value of data**

platform, analyzing, collection and visualization



Big Data Technology

- **Systems**

Your infrastructure must capitalize on real-time information flowing through your organization. It must be optimized for analytics to respond dynamically to the increasing demands of big data.

- **Privacy**


Your platform must comprise stringent policies and practices around privacy and insights on which your business relies.

- **Storage**


Your infrastructure must embody a defensible disposal strategy that reduces the run rate of storage and risk.

- **Security**

Your infrastructure must have strong security measures built in to guard your organization against internal and external threats.



Categories of Data Technology



Big Data Technology

Type	Consumer	Service Provided by Cloud	Service Level Coverage	Customization
SaaS	End user	<ul style="list-style-type: none"> Finished application 	<ul style="list-style-type: none"> Application uptime Application Performance 	<ul style="list-style-type: none"> Minimal to no customization Capabilities dictated by market or provider
PaaS	Application owner	<ul style="list-style-type: none"> Runtime environment for application code Cloud storage Other Cloud services such as integration 	<ul style="list-style-type: none"> Environment availability Environment performance No application coverage 	<ul style="list-style-type: none"> High degree of application level customization available within constraints of the service offered Many applications will need to be rewritten
IaaS	Application owner or IT provides OS, middleware, and application support	<ul style="list-style-type: none"> Virtual server Cloud storage 	<ul style="list-style-type: none"> Virtual server availability Time to provision No platform or application coverage 	<ul style="list-style-type: none"> Minimal constraints on applications installed on standardized virtual OS builds

Value/Opportunity of Big Data

- Discover the new role of data scientist

4.4 million data and analytics jobs is needed globally. only one-third of those jobs will be filled.

- Privacy, security and governance

While big data can provide significant value, it also presents significant risk.

- Create new business models with big data

Data-driven marketing and ad targeting and something new

4.4MILLION
data scientists
needed by 2015



Agenda

- Big Data
- Hadoop Brief Introduction
- HDFS Structure and Characteristic
- MapReduce Model
- Q&A

Three Papers from Google


December, 2004

Jeffrey Dean
Sanjay Ghemawat



MapReduce
*Simplified
Data Processing on
Large Clusters*

BigTable
*A Distributed
Storage System for
Structured Data*



November, 2006

Fay W. Chang
Jeffrey Dean
Sanjay Ghemawat
Wilson C. Hsie
etc.

**The Google
File System**

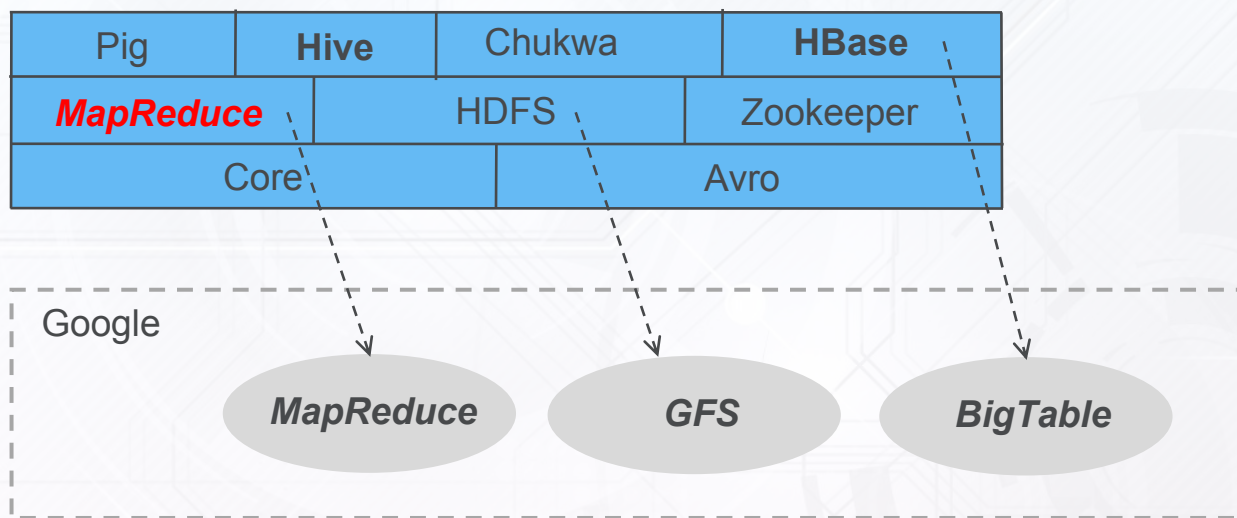
October, 2003

Sanjay Ghemawat
Howard Gobioff
Shun-Tak Leung



Hadoop Brief Introduction

- The Apache Hadoop project develops open-source software for reliable, scalable, distributed computing. Implemented in Java.
- The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. It is designed to scale up from single servers to thousands of machines, each offering local computation and storage. Rather than rely on hardware to deliver high-availability, the library itself is designed to detect and handle failures at the application layer, so delivering a highly-available service on top of a cluster of computers, each of which may be prone to failures.



Doug Cutting

The Companies around Hadoop



The Companies around Hadoop

- **Yahoo:**

More than 100,000 CPUs in more than 25,000 computers running it.

The biggest cluster: 4000 nodes, 2 x 4CPU boxes, with 4 x 1 TB disk, and 16 GB RAM.

- **Amazon :**

Process millions of sessions daily for analytics.

Using both Java and streaming APIs.

- **Facebook:**

Use it to store copies of internal log and dimension data sources.

a source for reporting and analytics, with machine learning algorithms.

- **BaiDu**

~10K cluster size

- **Sina**

>1K cluster size

.....

Agenda

- Big Data
- Hadoop Brief Introduction
- HDFS Structure and Characteristic
- MapReduce Model
- Q&A

HDFS Brief Introduction

- Hadoop Distributed File System (HDFS) is the primary storage system used by Hadoop applications. HDFS creates multiple replicas of data blocks and distributes them on compute nodes throughout a cluster to enable reliable, extremely rapid computations.
- For huge data (PB)
- Large scale distributed file system
- Cheap PC servers
- Based on file system of OS
- For rarely modified data
- Reliable

HDFS Brief Introduction

- **Master-Slave architecture**

Master: Name node (1, single-point failure)

Slave: Data node (N) SecondaryNameNode(1)

- **Distributed**

Every long file is split into blocks of 64MB, each block is allocated to different storage node.

- **Checkpoint:**

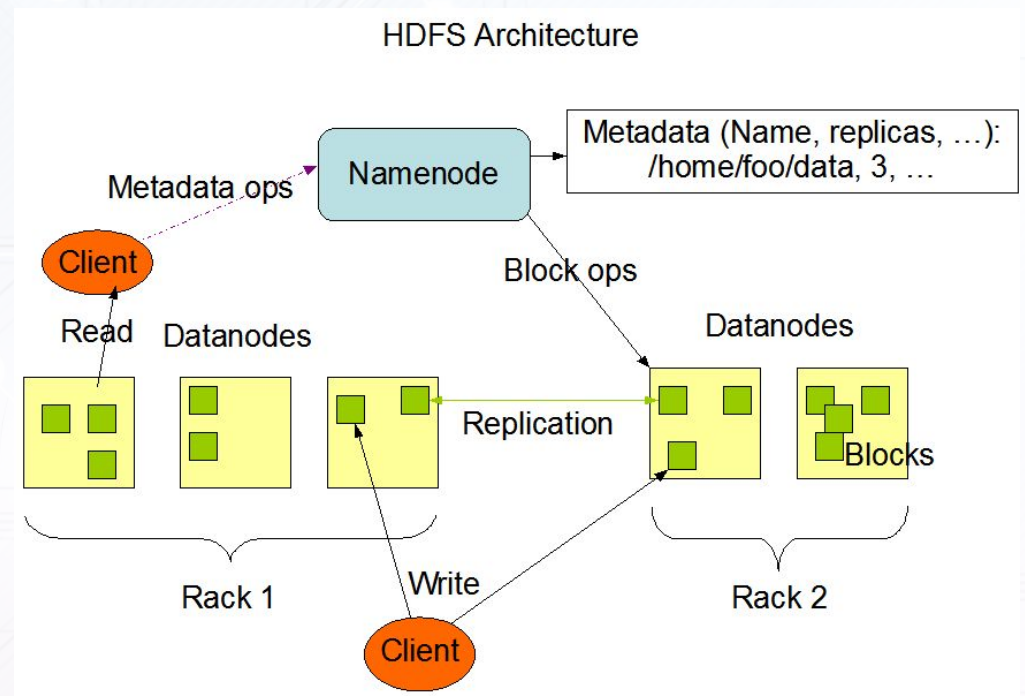
fsimage + edit log

- **Reliability**

Multiple replicas for each block
(default: 3)

- **Scalable**

Add disks/nodes



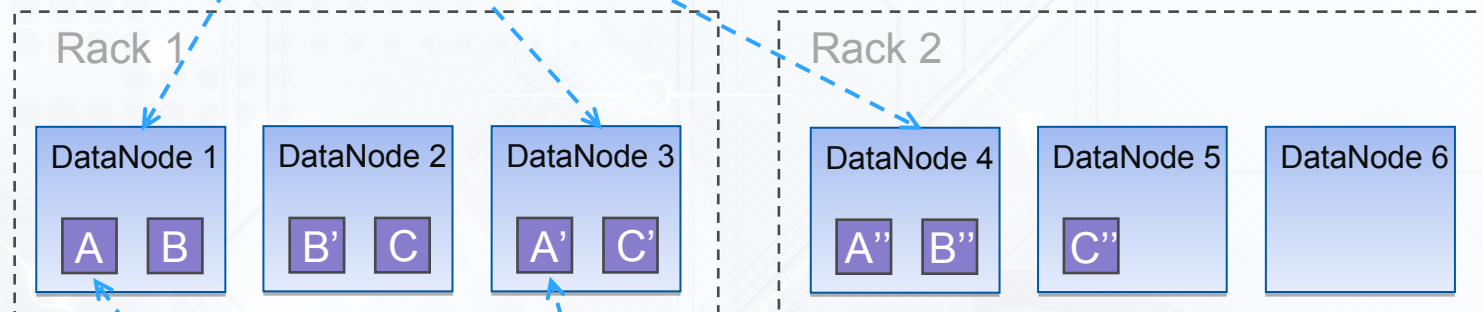
Upload file (140MB, /usr/file1)

Splitting:



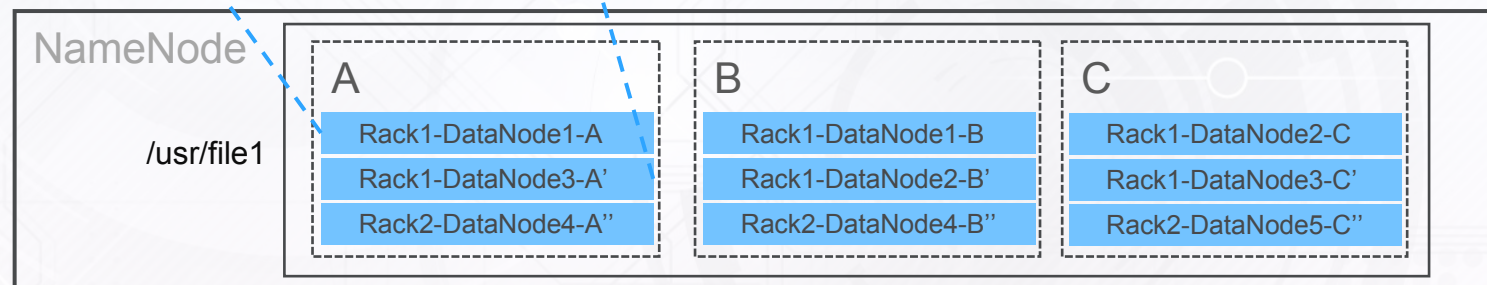
Replication:

HDFS writes the replica into different nodes in different racks. Enhance the reliability, but reduce the write speed.



Consistency:

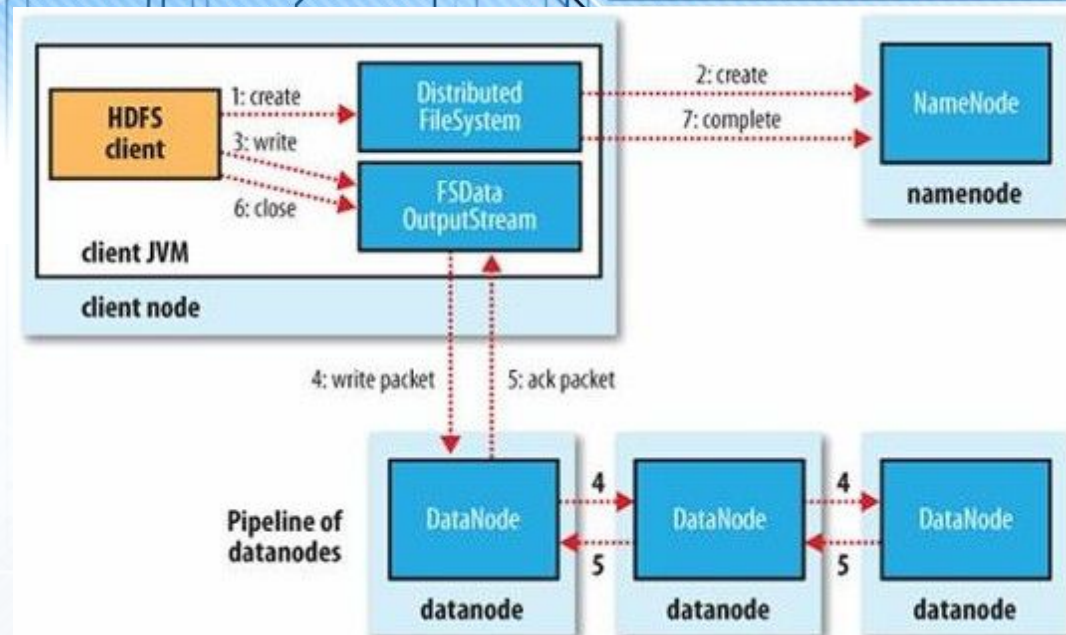
Read the nearest replica



HDFS create/write

FSDaOutputStream is returned by DistributedFileSystem after contacting NameNode.

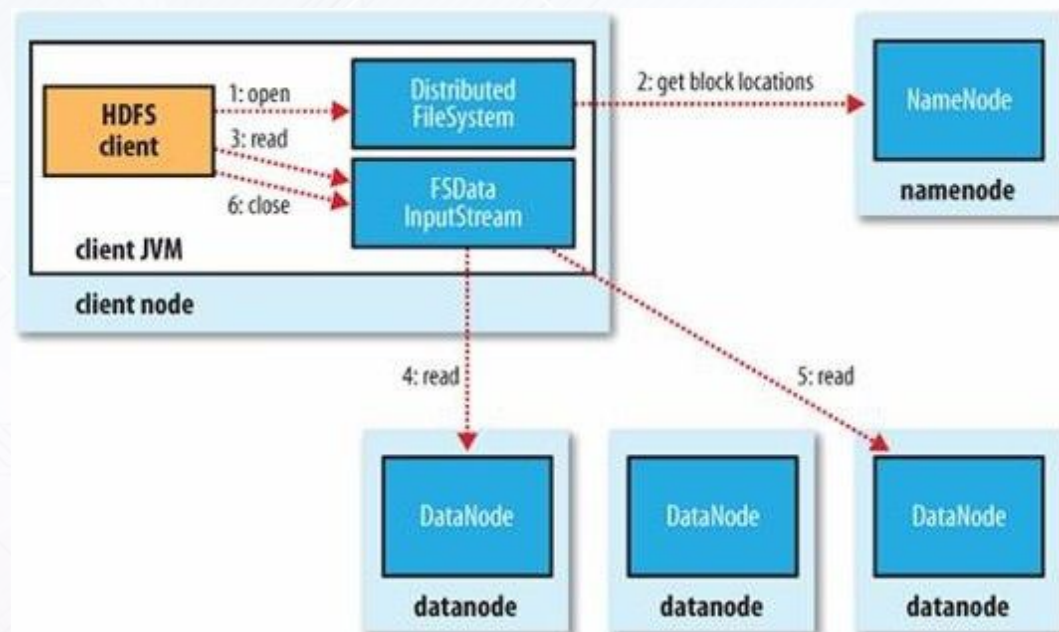
Upload file is split into multiple packets. Only after the **first** block has been stored into all the replica, the **second** block begins to upload.



HDFS read

FSDaOutputStream is returned by DistributedFileSystem after contacting NameNode to find the block address.

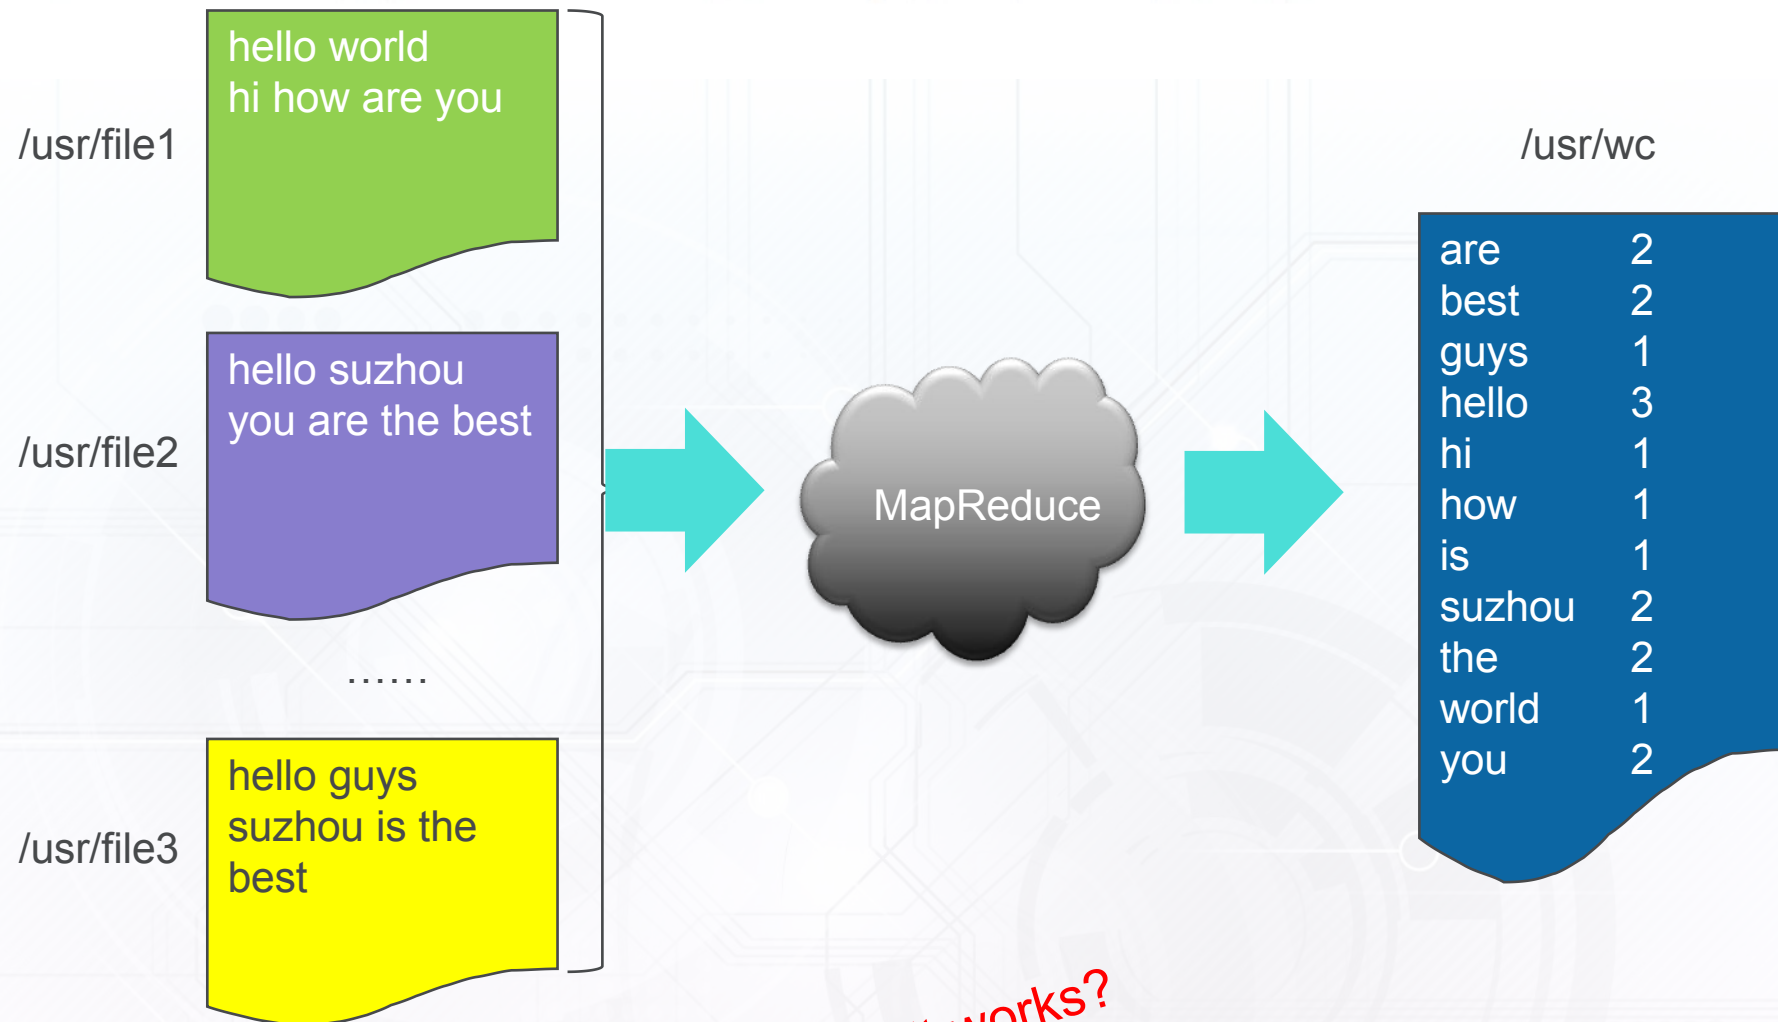
FSDaOutputStream read the block from the **nearest** datanode, if fails it goes to the replica.



Agenda

- Big Data
- Hadoop Brief Introduction
- HDFS Structure and Characteristic
- MapReduce Model
- Q&A

MapReduce Model



How it works?

MapReduce Model

- Two major functions – map and reduce

- **Map function**

Written by user

Takes input pairs and produce a set of intermediate key/value pairs

Intermediate key/value pairs are grouped and passed to reduce function

- **Reduce function**

Accepts an intermediate key and a set of values for the key

Merges together these values

Major components of Hadoop MapReduce 1.x

- **Job Tracker**

Cluster manager

Job manager

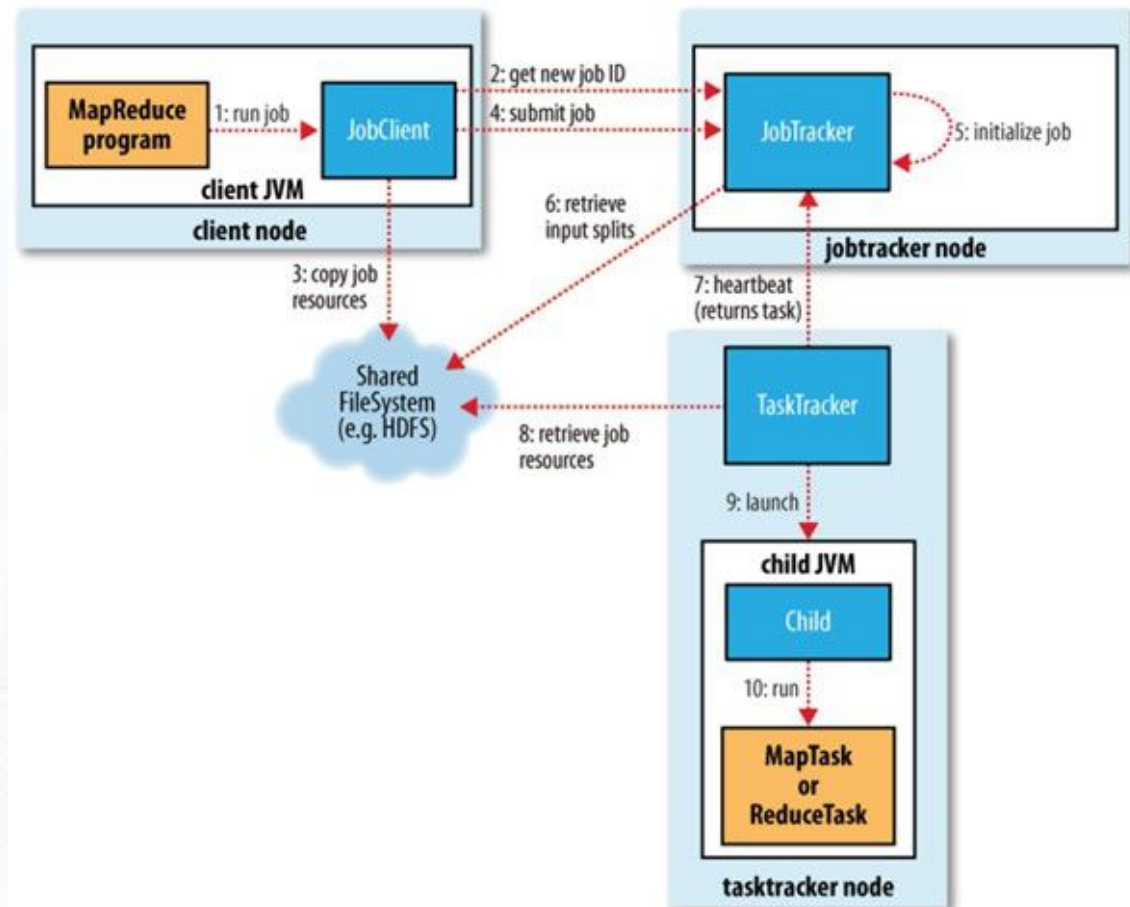
Job scheduler

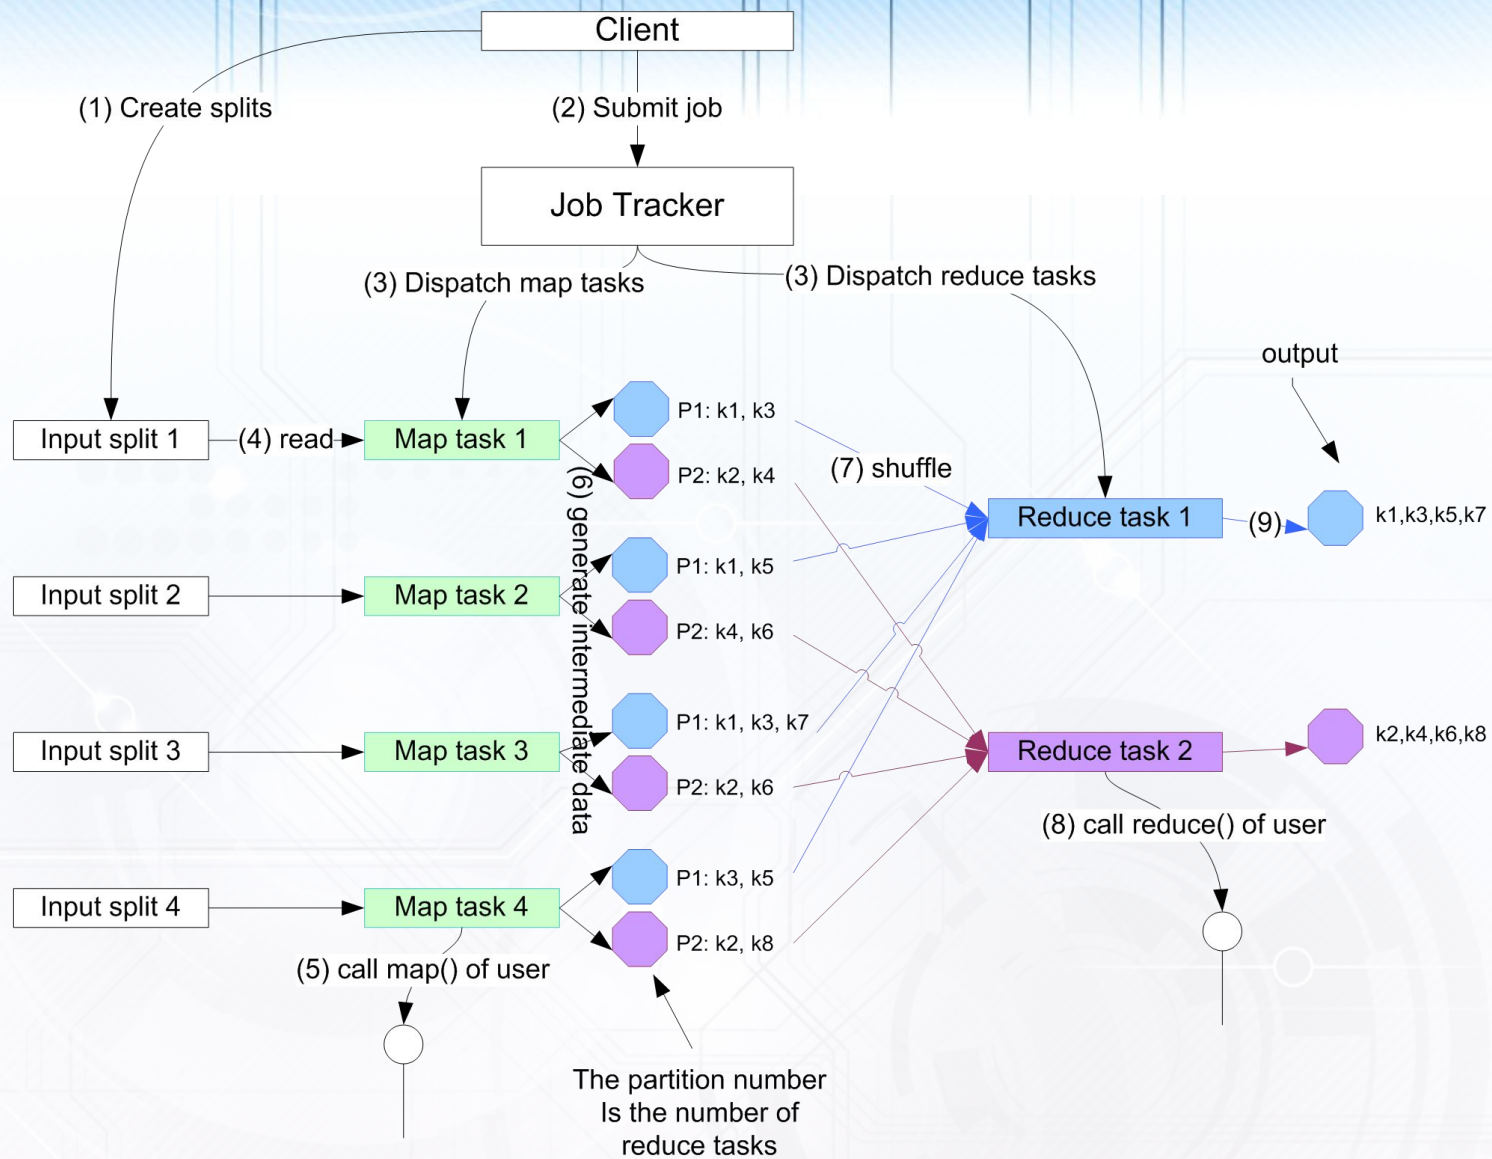
- **Task tracker**

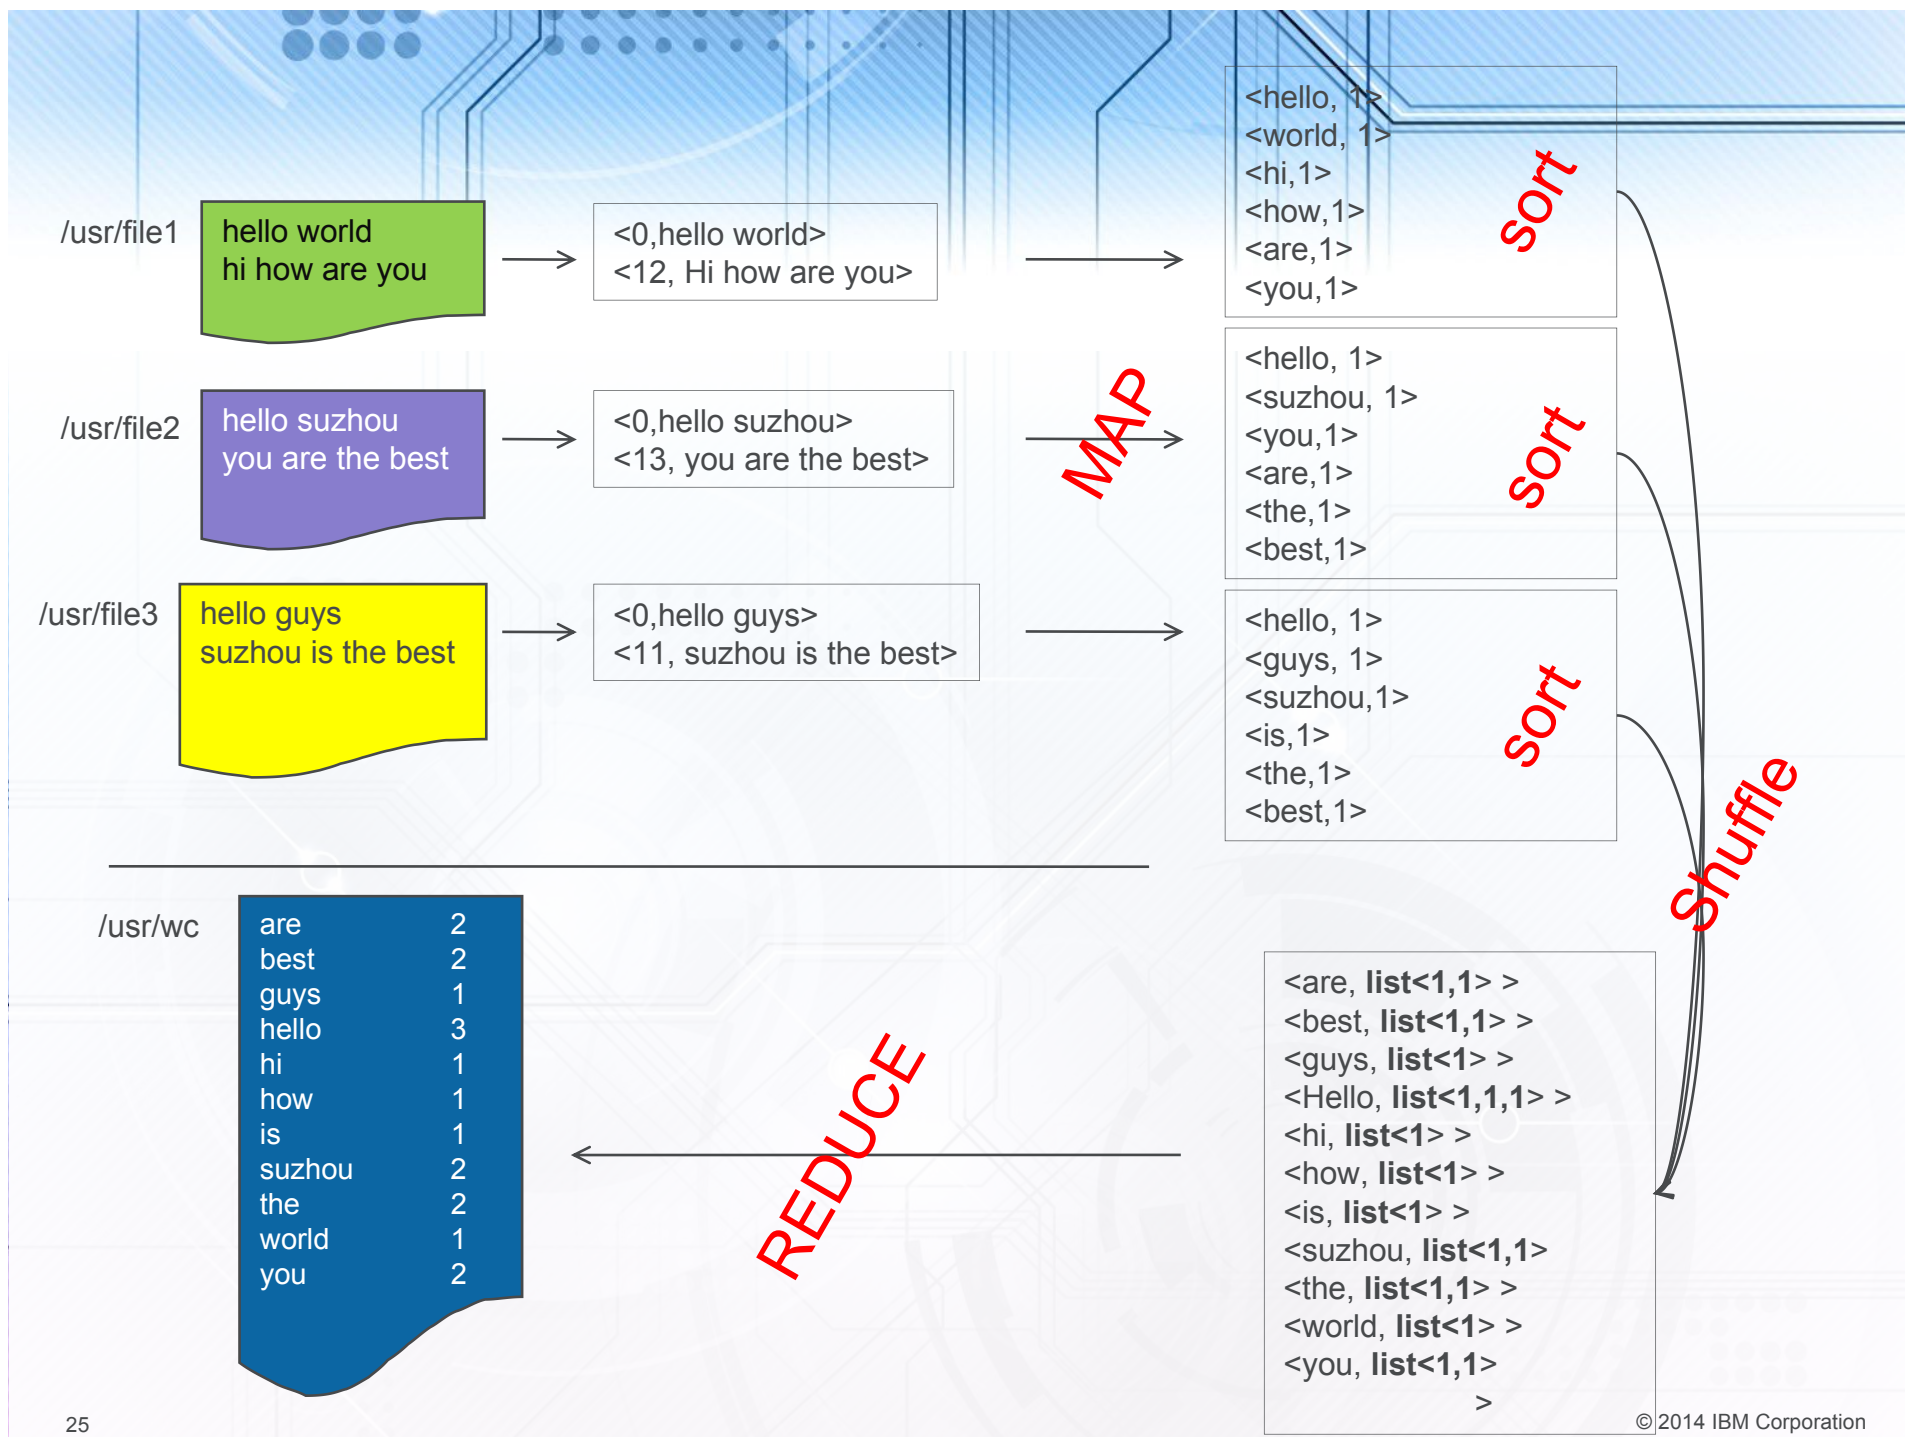
Task JVM spawner

Monitor

Shuffle server







Q&A

Q&A

参考资料

<http://hadoop.apache.org>